

# Project 1

## Generalized Linear Models 2020

Lene Tillerli Omdal LeneTOM@stud.ntnu.no 10022

Arne Rustad ArneIR@stud.ntnu.no 10028

02 oktober, 2020

1)

a)

Each observation  $y_i$  is Poisson distributed with expectation  $\lambda_i$ ,  $y_i \sim \text{poisson}(\lambda_i)$ . The probability density function of  $y_i$  is

$$f(y_i|\lambda_i) = \frac{\lambda_i^{y_i}}{y_i!} e^{-\lambda_i}.$$

Using the canonical link function we get the following relationship between  $E[y_i] = \mu_i$  and the linear predictor  $x_i^T \beta = \eta_i$

$$\ln(\mu_i) = \ln(\lambda_i) = x_i^T \beta.$$

This gives the likelihood function

$$\begin{aligned} L(\beta|y_i, x_i) &= \prod_{i=1}^n f(y_i|\lambda_i) = \prod_{i=1}^n f(y_i|\beta, x_i) \\ &= \prod_{i=1}^n \frac{(e^{x_i^T \beta})^{y_i}}{y_i!} e^{-e^{x_i^T \beta}}, \end{aligned}$$

and the log likelihood function

$$\begin{aligned} l(\beta|y_i, x_i) &= \ln(L(\beta|y_i, x_i)) = \sum_{i=1}^n \ln(f(y_i|\lambda_i)) \\ &= \sum_{i=1}^n -\ln(y_i!) + y_i x_i^T \beta - e^{x_i^T \beta}. \end{aligned}$$

The score vector becomes

$$\begin{aligned}
s(\beta) &= \frac{\partial l(\beta)}{\partial \beta} = \sum_{i=1}^n y_i x_i - e^{x_i^T \beta} \cdot \frac{\partial}{\partial \beta} x_i^T \beta \\
&= \sum_{i=1}^n (y_i - e^{x_i^T \beta}) x_i \\
&= \sum_{i=1}^n (y_i - \lambda_i) x_i.
\end{aligned}$$

In matrix notation this is

$$s(\beta) = X^T (Y - \lambda),$$

where

$$X = \begin{bmatrix} x_1^T \\ x_2^T \\ \vdots \\ x_n^T \end{bmatrix} = \begin{bmatrix} x_{11} & x_{12} & \dots & x_{1p} \\ x_{21} & x_{22} & \dots & x_{2p} \\ \vdots & \vdots & \ddots & \vdots \\ x_{n1} & x_{n2} & \dots & x_{np} \end{bmatrix}$$

and

$$\lambda = \begin{bmatrix} \lambda_1 \\ \lambda_2 \\ \vdots \\ \lambda_n \end{bmatrix} = \begin{bmatrix} e^{x_1^T \beta} \\ e^{x_2^T \beta} \\ \vdots \\ e^{x_n^T \beta} \end{bmatrix}.$$

The observed Fisher information is given as

$$H(\beta) = -\frac{\partial^2 l(\beta)}{\partial \beta \partial \beta^T} = -\frac{\partial s(\beta)}{\partial \beta^T}.$$

Inserting the obtained score function results in

$$\begin{aligned}
H(\beta) &= -\frac{\partial}{\partial \beta^T} \sum_{i=1}^n (y_i - e^{x_i^T \beta}) x_i \\
&= \sum_{i=1}^n e^{x_i^T \beta} x_i \frac{\partial}{\partial \beta^T} (x_i^T \beta) \\
&= \sum_{i=1}^n e^{x_i^T \beta} x_i x_i^T,
\end{aligned}$$

given in matrix notation this is

$$H(\beta) = X^T \text{diag}(\lambda) X.$$

The expected Fisher information then becomes

$$\begin{aligned} F(\beta) &= E[H(\beta)] \\ &= \sum_{i=1}^n e^{x_i^T \beta} x_i x_i^T = X^T \text{diag}(\lambda) X \\ &= H(\beta). \end{aligned}$$

b)

```
vec2diag = function(vec) {
  m = diag(length(vec))
  for (i in 1:length(vec)) {
    m[i, i] = vec[i]
  }
  return(m)
}

diag2vec = function(m) {
  vec = rep(0, nrow(m))
  for (i in 1:nrow(m)) {
    vec[i] = m[i, i]
  }
  return(vec)
}

myglm = function(formula, data, start = rep(0, ncol(model.matrix(formula, data))),
  epsilon = 1e-08) {
  n = nrow(data)
  X = model.matrix(formula, data)
  Y = matrix(data$y)
  beta = matrix(start)
  delta.beta = epsilon + 1
  counter = 0
  while (Norm(delta.beta) > epsilon) {
    lambda = matrix(exp(X %*% beta))
    F.beta = t(X) %*% vec2diag(lambda) %*% X
    s.beta = t(X) %*% (Y - lambda)
    delta.beta = solve(F.beta, s.beta)
    beta = beta + delta.beta
    counter = counter + 1
  }
  logY = ifelse(Y == 0, 1, log(Y))
  deviance = 2 * sum(Y * logY - Y - Y * log(lambda) + lambda)
  vcov = solve(F.beta)
  std.errors = sqrt(diag2vec(vcov))
  Coefficients = matrix(c(beta, std.errors), nrow = length(beta))
  colnames(Coefficients) = c("Estimate", "Std. Error")
  rownames(Coefficients) = rownames(beta)
  return(list(Coefficients = Coefficients, Deviance = deviance, vcov = vcov,
```

```

      X = X, p = length(beta), n = length(Y)))
}

```

c)

We test our function by comparing the results it gives on simulated data, to those obtained using the built-in functions `glm` and `vcov`.

```

n = 1000
real_beta1 = 0.1
real_beta0 = 1
t = runif(n, 0, 10)
y = rpois(n, exp(real_beta0 + real_beta1 * t))
test.data = data.frame(y = y, t = t)
mymodel = myglm(y ~ t, test.data)
model_compare = glm(y ~ t, family = "poisson", data = test.data)

```

Comparing the coefficients and std. errors:

```
mymodel$Coefficients
```

```

##              Estimate Std. Error
## (Intercept) 1.00870950 0.034038918
## t           0.09831261 0.005118018

```

```
summary(model_compare)$coefficients
```

```

##              Estimate Std. Error z value    Pr(>|z|)
## (Intercept) 1.00870951 0.034038001 29.63480 5.324056e-193
## t           0.09831261 0.005117925 19.20947 3.084176e-82

```

Comparing the deviance:

```
mymodel$Deviance
```

```
## [1] 1034.105
```

```
model_compare$deviance
```

```
## [1] 1034.105
```

Comparing the covariance matrix of the coefficients:

```
mymodel$vcov
```

```

##              (Intercept)          t
## (Intercept)  0.001158648 -1.576250e-04
## t           -0.000157625  2.619411e-05

```

```
vcov(model_compare)
```

```
##              (Intercept)              t
## (Intercept)  0.0011585855 -1.576173e-04
## t           -0.0001576173  2.619316e-05
```

All of the elements appear to be computed correctly.

2)

a)

We interpret the parameter  $\lambda_0$  as the expected number of fledglings produced by each female if she starts breeding at the optimal time.  $\theta$  we interpret as the optimal timing for breeding, based on the timing of important food resources. Lastly, we interpret  $\omega$  as a parameter for how fast the expected number of fledglings produced by each female deteriorates when initiating breeding at a non-optimal date. If  $\omega$  is small then  $\lambda_i$  decreases rapidly once  $t_i$  differs from  $\theta$ . If, however,  $\omega$  is large, a difference between  $t_i$  and  $\theta$  is less severe concerning the decrease of  $\lambda_i$ .

b)

The poisson distribution belongs to the exponential family and we choose to use the canonical link function.

$$\begin{aligned}\mu_i &= \lambda_i = \lambda_0 e^{\frac{(t_i - \theta)^2}{2\omega^2}} \\ \ln(\mu_i) &= \ln(\lambda_0) - \frac{(t_i - \theta)^2}{2\omega^2} \\ \ln(\mu_i) &= \ln(\lambda_0) - \frac{\theta^2}{2\omega^2} + \frac{\theta}{\omega^2} t_i - \frac{1}{2\omega^2} t_i^2 \\ g(\mu_i) &= t_i^T \beta\end{aligned}$$

After reparametrization the expectation of  $\mu_i$  is connected to a set of linear predictors through a link function, therefore this is a GLM. The relationship between the GLM-parameters  $\beta$  and  $\omega$ ,  $\theta$  and  $\lambda_0$  is defined by

$$t_i^T \beta = \ln(\lambda_0) - \frac{\theta^2}{2\omega^2} + \frac{\theta}{\omega^2} t_i - \frac{1}{2\omega^2} t_i^2.$$

This gives  $\beta_0 = \ln(\lambda_0) - \frac{\theta^2}{2\omega^2}$ ,  $\beta_1 = \frac{\theta}{\omega^2}$  and  $\beta_2 = \frac{-1}{2\omega^2}$ .

c)

```
load(url("https://www.math.ntnu.no/emner/TMA4315/2020h/hoge-veluwe.Rdata"))
df.birds = data
model = myglm(y ~ t + I(t^2), data = df.birds)
```

d)

```

model.wo.t2 = myglm(y ~ t, data = df.birds)

compare.nested.models = function(model0, model1) {
  LRT = model0$Deviance - model1$Deviance
  r = model1$p - model0$p
  probLRT = pchisq(LRT, df = r, lower.tail = FALSE)
  if (probLRT == 0)
    probLRT = "<1e^-16"
  m = matrix(c(LRT, probLRT), nrow = 1)
  colnames(m) = c("LRT", "Pr(>Chi)")
  return(m)
}

lrt.test = compare.nested.models(model.wo.t2, model)
lrt.test

```

```

##          LRT      Pr(>Chi)
## [1,] 11.92985 0.0005524157

```

We perform a hypothesis test where the null hypothesis is that there is no quadratic effect of  $t$ , against the alternative hypothesis that there exist a non-neglectable quadratic effect of  $t$ . The likelihood-ratio test gives that if  $H_0$  is correct it is only a  $5.5241571 \times 10^{-4}$  probability to observe equally or more extreme results. Based on this we reject the null hypothesis and conclude that there is evidence of a quadratic effect of  $t$ .

e)

```

goodness.of.fit = function(model) {
  prob = pchisq(model$Deviance, model$n - model$p, lower.tail = FALSE)
  m = matrix(c(model$Deviance, prob), ncol = 2)
  colnames(m) = c("Deviance", "Pr(>Chi)")
  rownames(m) = "Model"
  return(m)
}

goodness.of.fit(model)

```

```

##          Deviance      Pr(>Chi)
## Model 277.4613 2.210642e-12


```

The model fits the data quite well, at least according to the goodness of fit test using deviance shown above. Plotting a histogram with  $y_i$  as parameter we get

```

ggplot(df.birds, aes(y)) + geom_bar()

```



It is difficult to assess from the plot above if the model assumptions are violated, since the histogram shows observations from a mix of different poisson models (different expectation and variance). Therefore to conclude we take a closer look at the distribution of  $\lambda_i$ -values.

```
beta = model$Coefficients[, 1]
X = model$X
lambda = exp(X %*% beta)
lambda.min = min(lambda)
numb.lambdas.less.than.4 = length(lambda[lambda <= 4])
```

The smallest  $\lambda$  value is  $\lambda_{min} = 1.89$ . Additionally, out of 135 observations, we only estimate 5 observations to be drawn from a poisson distribution with expectation less than or equal to 4. Taking this into account, it is highly unlikely to observe 14 cases of  $y_i = 0$  if the assumption of each  $y_i$  being poisson distributed is satisfied. However, even though the assumption of each  $y_i$  being poisson distributed might be violated, it still appears that the model has good overall predictive ability. The model just underestimates the chance of  $y_i = 0$ .

To prove the claim that it is highly unlikely that each  $y_i$  is poisson distributed with expectation  $\lambda_i$ , we want to find an upper limit for the probability of observing 14 or more cases of  $y_i = 0$  given that  $y_i$  is poisson distributed. To do this we divide the observations into two groups. The first containing the  $y_i$ s with  $\lambda_i \leq 4$ , and the second containing the ones with  $\lambda_i > 4$ . Since we only need an upper limit we assume all five observations in the first group was  $y_i = 0$ . Then we can study the probability of the second group getting  $14 - 5 = 9$  observations equal to zero. Again, since we only need an upper bound, we assume each observation in second group has  $\lambda_i = 4$ . Then the probability of at least 9  $y_i = 0$  among the 130 observations in the second group, can be written as

$$\sum_{i=9}^{130} \frac{130}{i! \cdot (130-i)!} e^{-4i} (1 - e^{-4})^{130-i}.$$

```
p = 0
for (i in 9:130) {
  k = factorial(130)/(factorial(i) * factorial(130 - i)) * exp(-4 * i) * (1 -
    exp(-4 * (130 - i)))
  p = p + k
}
```

Calculating this probability we get an upper bound for the probability of observing something at least as extreme as we have done, given the assumption of  $y_i \sim \text{poisson}(\lambda_i)$ . This upper bound is 0.0065138. Consequently, it is highly likely that our assumption of poisson distribution is violated.

f)

We can now easily find the maximum likelihood estimates of  $\theta$  and  $\omega$  since we have the maximum likelihood estimates of the  $\beta$  coefficients and we have three equations describing the relationship between  $\lambda_0$ ,  $\theta$ ,  $\omega$  and  $\beta$ . Solving two of these equations gives expressions for  $\hat{\theta}$  and  $\hat{\omega}$ .

$$\hat{\omega}(\hat{\beta}) = \sqrt{\frac{-1}{2\hat{\beta}_2}}$$

$$\hat{\theta}(\hat{\beta}) = \hat{\beta}_1 \hat{\omega}^2 = \frac{-\hat{\beta}_1}{2\hat{\beta}_2}$$

```
omega.hat = sqrt(-1/(2 * beta[3]))
theta.hat = beta[2] * omega.hat^2
```

The estimates are  $\hat{\omega} = 12.312$  and  $\hat{\theta} = 12.912$ .

Using the Delta method we can find the variance of  $\hat{\omega}(\hat{\beta})$  and  $\hat{\theta}(\hat{\beta})$ . Let  $G$  be a function of  $\hat{\beta}$  then

$$\text{Var}[G(\hat{\beta})] = \nabla G(\hat{\beta})^T \text{Cov}[\hat{\beta}] \nabla G(\hat{\beta}). \quad (1)$$

The gradient of  $\hat{\omega}(\hat{\beta})$  is

$$\nabla \hat{\omega}(\hat{\beta})^T = [0 \quad 0 \quad (-2\hat{\beta}_2)^{-3/2}].$$

The gradient of  $\hat{\theta}(\hat{\beta})$  is

$$\nabla \hat{\theta}(\hat{\beta})^T = \left[ 0 \quad \frac{-1}{2\hat{\beta}_2} \quad \frac{\hat{\beta}_1}{2\hat{\beta}_2^2} \right].$$

Inserting these gradients into (1) we compute the variance of  $\hat{\omega}$  and  $\hat{\theta}$

```
grad.omega = matrix(c(0, 0, (-2 * beta[3])^(-3/2)))
se.omega.hat = sqrt(t(grad.omega) %*% model$vcov %*% grad.omega)

grad.theta = matrix(c(0, -1/(2 * beta[3]), beta[2]/(2 * beta[3]^2)))
se.theta.hat = sqrt(t(grad.theta) %*% model$vcov %*% grad.theta)

coef = matrix(c(c(omega.hat, theta.hat), c(se.omega.hat, se.theta.hat)), ncol = 2)
colnames(coef) = c("Estimate", "Std. Error")
rownames(coef) = c("Omega", "Theta")
coef
```

```
##      Estimate Std. Error
## Omega 12.31175    1.902526
## Theta 12.91197    1.609386
```

Using the Delta method we get  $\text{SE}[\hat{\omega}] = 1.903$  and  $\text{SE}[\hat{\theta}] = 1.609$ .



g)

We want to examine if the optimal breeding date,  $\hat{\theta}$ , differs from the mean breeding date of the population,  $\mu$ . In task 2f we estimate the optimal breeding date as  $\hat{\theta} = 12.91$ . An estimator for  $\mu$  is

$$\hat{\mu} = \sum_{i=1}^n t_i.$$

```
mu.hat = mean(df.birds$t)
```

An estimate for  $\mu$  is therefore  $\hat{\mu} = 15.937037$ . To decide if the mean value of  $t$  in this population is significantly different from the estimated optimal date we perform a hypothesis test. Choosing  $\alpha = 0.05$  as our significance level.

$$\begin{aligned} H_0 : & \quad \mu = \theta \\ H_1 : & \quad \mu \neq \theta \end{aligned}$$

We assume that each  $t_i$  is identically and independently distributed. Then we have by the central limit theorem that asymptotically  $\hat{\mu}$  becomes normally distributed as  $n \rightarrow \infty$ . We also know that  $\hat{\theta}$  is asymptotically normally distributed. Consequently  $\hat{\mu} - \hat{\theta}$  is asymptotically normally distributed, since it is a linear combination of asymptotically normally distributed variables. From the definition of  $\hat{\mu}$  it is clear that it is an unbiased estimator. We assume  $\hat{\theta}$  is an unbiased estimator, and that  $\hat{\theta}$  and  $\hat{\mu}$  is independent of each other. This gives

$$\begin{aligned} E[\hat{\mu} - \hat{\theta}] &= \mu - \theta \\ \text{Var}[\hat{\mu} - \hat{\theta}] &= \text{Var}[\hat{\mu}] + \text{Var}[\hat{\theta}]. \end{aligned}$$

Since  $\hat{\mu} - \hat{\theta}$  is asymptotically normally distributed and number of observations,  $n = 135$ , is quite large we have that

$$Z = \frac{(\hat{\mu} - \hat{\theta}) - (\mu - \theta)}{\text{SE}[\hat{\mu} - \hat{\theta}]} \approx \frac{(\hat{\mu} - \hat{\theta}) - (\mu - \theta)}{\widehat{\text{SE}}[\hat{\mu} - \hat{\theta}]}$$

is approximately standard normal distributed.

$$\begin{aligned} P(|Z| > z_\alpha) &= P(-z_\alpha \leq \frac{(\hat{\mu} - \hat{\theta}) - (\mu - \theta)}{\text{SE}[\hat{\mu} - \hat{\theta}]} \leq z_\alpha) \\ &= P((\hat{\mu} - \hat{\theta}) - z_\alpha \text{SE}[\hat{\mu} - \hat{\theta}] \leq \mu - \theta \leq (\hat{\mu} - \hat{\theta}) + z_\alpha \text{SE}[\hat{\mu} - \hat{\theta}]) \\ &\approx P((\hat{\mu} - \hat{\theta}) - z_\alpha \widehat{\text{SE}}[\hat{\mu} - \hat{\theta}] \leq \mu - \theta \leq (\hat{\mu} - \hat{\theta}) + z_\alpha \widehat{\text{SE}}[\hat{\mu} - \hat{\theta}]) \end{aligned}$$

```
var.mu.hat = var(df.birds$t)/length(df.birds$t)
z.value = (mu.hat - theta.hat)/sqrt(var.mu.hat + se.theta.hat^2)
p.value = 2 * pnorm(z.value, lower.tail = F)
result.z.test = matrix(c(z.value, p.value), nrow = 1)
colnames(result.z.test) = c("Z-value", "Pr(>|Z|)")
rownames(result.z.test) = "H0: mu = theta"
result.z.test
```

```
##                Z-value    Pr(>|Z|)
## H0: mu = theta 1.821029 0.06860242
```

Since the p-value is larger than the significance level  $\alpha = 0.05$  we choose to not reject the null hypothesis. In other words we conclude that the mean breeding date for the population is not significantly different from the optimal breeding date.

3)

```
B = 1000
n = length(df.birds$y)
curr_data = df.birds
formel = formula(y ~ t + I(t^2))
X = model.matrix(formel, curr_data)
lambda = matrix(exp(X %*% beta))
betas = matrix(, nrow = length(beta), ncol = B)
for (i in 1:B) {
  y.sim = rpois(n, lambda)
  curr_data$y = y.sim
  fit = myglm(formel, curr_data, start = beta)
  betas[, i] = matrix(fit$Coefficients[, 1])
}

df.betas.bootstrap = data.frame(beta0 = betas[1, ], beta1 = betas[2, ], beta2 = betas[3,
  ])

var(df.betas.bootstrap)
```

```
##                beta0                beta1                beta2
## beta0  0.0810534429 -9.375393e-03  2.543442e-04
## beta1 -0.0093753927  1.160409e-03 -3.356127e-05
## beta2  0.0002543442 -3.356127e-05  1.038142e-06
```

```
model$vcov
```

```
##                (Intercept)                t                I(t^2)
## (Intercept)  0.0797693783 -9.308596e-03  2.550195e-04
## t            -0.0093085957  1.159672e-03 -3.369024e-05
## I(t^2)       0.0002550195 -3.369024e-05  1.039306e-06
```

The results from parametric bootstrapping are quite similar to those obtained from the expected Fisher information. This indicates that the asymptotic results is a good approximation even for finite sample sizes ( $n=135$ ).